
NAND Flash & SSDs

Angelos Bilas

FORTH-ICS

and University of Crete, Greece

bilas@ics.forth.gr

Material

- Slides
 - Toshiba slides (first NAND in 1989)
 - Al Fazio, Intel Fellow, Director, Memory technology Development, November 12, 2008
 - Reading
 - (SSDs) Nitin Agrawal et al. Design tradeoffs for SSD performance. In USENIX 2008 Annual Technical Conference (ATC'08).
 - (NVM Intro) Natarajan, S. et al., "Searching for the dream embedded memory," Solid-State Circuits Magazine, IEEE , vol.1, no.3, pp.34,44, Summer 2009.
 - (Flash Intro) Micheloni, R. et al., "Non-Volatile Memories for Removable Media," Proceedings of the IEEE , vol.97, no.1, pp.148,160, Jan. 2009.
 - (edited book) Inside Solid State Drives (SSDs) by R. Micheloni et al., Springer Verlag.
-

Outline

- NAND Flash
 - How it works
 - Structure
 - Errors
 - SSDs
 - Structure
 - Garbage collection
 - Wear leveling
 - System-level SSD Storage
 - SSD caching
 - Remarks
 - Non-volatile memories (NVM)
-


20+ Years Flash Floating Gate Technology



1986/1.5µm




1988/1.0µm




1991/0.8µm



1993/0.6µm



1996/0.4µm



1998/0.25µm



2000/0.18µm



2002/0.13µm



2004/90nm



2006/72nm



2007/50nm



2008/34nm
32Gbit

*3+ decades of floating-gate technology
scaling starting on EPROM → Flash*

*Technology designed for high-volume
manufacturing*

1987 View of NVM in Compute



A Full Range of NVM in Compute



Intel® X25-E Extreme
SATA Solid-State Drive




Intel® X25-M and X18-M
Mainstream SATA Solid-State Drive



Intel®
Turbo Memory



Intel® Z-P230 PATA
Solid-State Drive



Intel® Z-P140 PATA
Solid-State Drive



Intel® Z-U130 USB
Solid-State Drive




Source: Intel



	NAND	NOR
Cell Array	<p>Word line</p> <p>Unit Cell</p> <p>Source line</p>	<p>Bit line</p> <p>Word line</p> <p>Contact</p> <p>Unit Cell</p> <p>Source line</p>
Layout	<p>$2F$</p> <p>$2F$</p>	<p>$2F$</p> <p>$5F$</p>
Cross-section		
Cell size	$4F^2$	$10F^2$


Flash Cell Layout and Cross-Section



N-Channel MOSFET with a few distinguishing features:

- Isolated floating gate
- Charge storage on Floating gate modulates threshold voltage of underlying MOSFET

Charge Storage: Program and Erase




Programming: NAND


Programming means injecting electrons to the FG

- Fowler-Nordheim Tunneling

Erase: Fowler-Nordheim Tunneling in reverse direction




Erase



Reliability


- Three main types of problems
- 1. Oxide degradation over time (data retention)
 - No accesses – FLASH on the shelf
 - Oxide degradation leads to bit errors
 - Unused FLASH, typical retention time 10 years
 - BUT depends on how much the material has been used (erase/prog cycles)
- 2. Oxide degradation due to erase/write
 - Main problem so far
 - Between 10,000 and 100,000 erase cycles
 - Wear leveling used to compensate (will discuss)
- 3. Bit errors due to reading neighboring cells
 - Read disturbance
 - Until now smaller problem, but starting to be relevant
 - Depends on how much the material has been used (erase/prog cycles)
- Read/program/erase process is electrically complex
 - Includes various mechanisms
 - Each mechanism with its own sources of error
- (NOR does not suffer from these problems)

ex.256Mb NAND Flash Memory



256Mb NAND Flash
Page Size : 512+16 Bytes
Block Size : 16KBytes
of Blocks : 2048 Blocks

NAND Physical Organization



NAND Operations

- Read
 - Setup time to “reach” a page within block (sequential)
 - Read a full page at a time
 - NOR can be read “truly” randomly
 - Program (1 \rightarrow 0)
 - Setup time to “reach” a page within block (sequential)
 - Program a full page at a time, set independent cell bits to 0
 - NOR can be programmed randomly but each word separate (therefore a page takes longer to program)
 - Erase (0 \rightarrow 1)
 - Can only erase a full block (physical organization and method)
 - Different methods to erase NAND/NOR
-

	NAND	NOR
Capacity	~ 1Gbit (2chips/pkg)	~ 128Mbit
Power Supply	2.7-3.6V	2.3-3.6V
I/O	x8	x8/x16
Access Time	50ns(serial access cycle) 25μs(random access)	70ns(30pF, 2.3V) 65ns(30pF, 2.7V)
Program Speed (typ.)	—	8μs/Byte
	200μs/512Byte	4.1ms/512Byte
Erase Speed(typ.)	2ms/Block (16KB)	700ms/Block
Prog+Erase(typ.)	33.6ms / 64KB	1.23s/Block (main:64KB)


So What?

- Free lunch
 - Use NAND for large, sequential operations
 - Small, random READs are great too
 - Use NAND for workloads with few writes
 - Mobile is a good match: few updates, media files
 - When FLASH is good, it's very good!
 - Can we use it elsewhere?
 - General purpose computing and servers have an I/O problem
 - No free lunch there
 - Small random writes + many writes
 - Reliability
-

Potential

- A disk is about 100 IOPS and 100 Mbytes/s
- An SSD can do 100,000 IOPS and 1000 MBy/s
- Does it matter?
- NO: use many disks
- How many?
 - 100K IOPS/core
 - 100 MBy/s/core
- A server with 128 cores is not easy to feed from disks

Motivation for NVM in compute: Huge Scaling Discrepancy Between CPU and HDD



1.3X vs 175X in 13 years!



Flash in Storage Systems

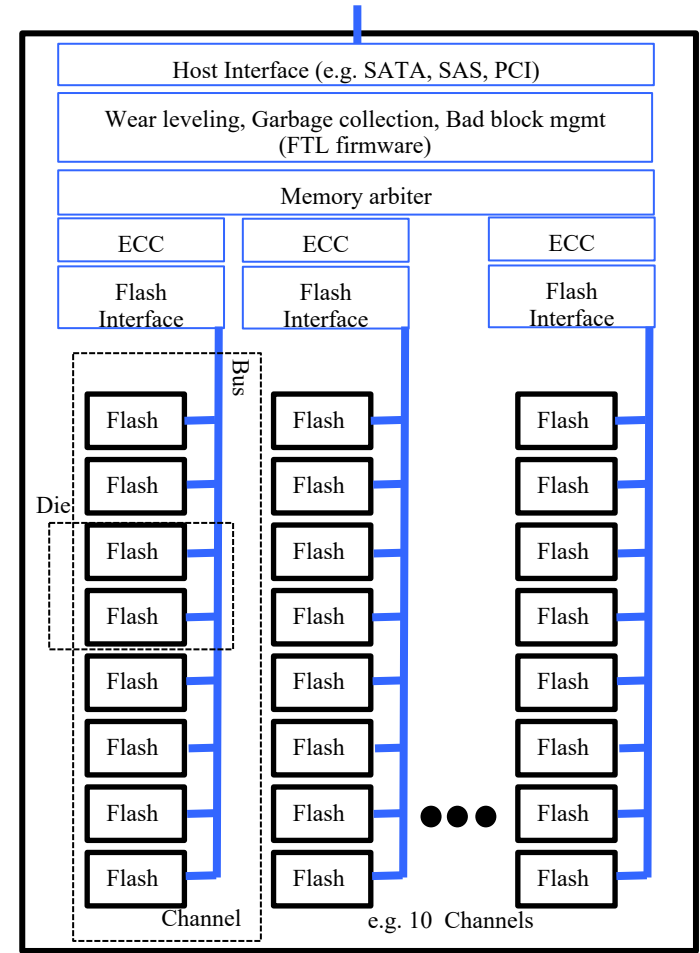
- How do we use Flash to improve storage I/O?
- Solid state drives (SSDs)

Outline

- NAND Flash
 - How it works
 - Structure
 - Errors
 - SSDs
 - I/O Concurrency
 - Garbage collection
 - Wear leveling
 - System-level SSD Storage
 - SSD caching
 - Remarks
 - Non-volatile memories (NVM)
-

SSD Organization

- 10 channels
- 2 components per channel
- 4 dies per component
- Multiple outstanding commands per channel
- Error detection/correction per channel



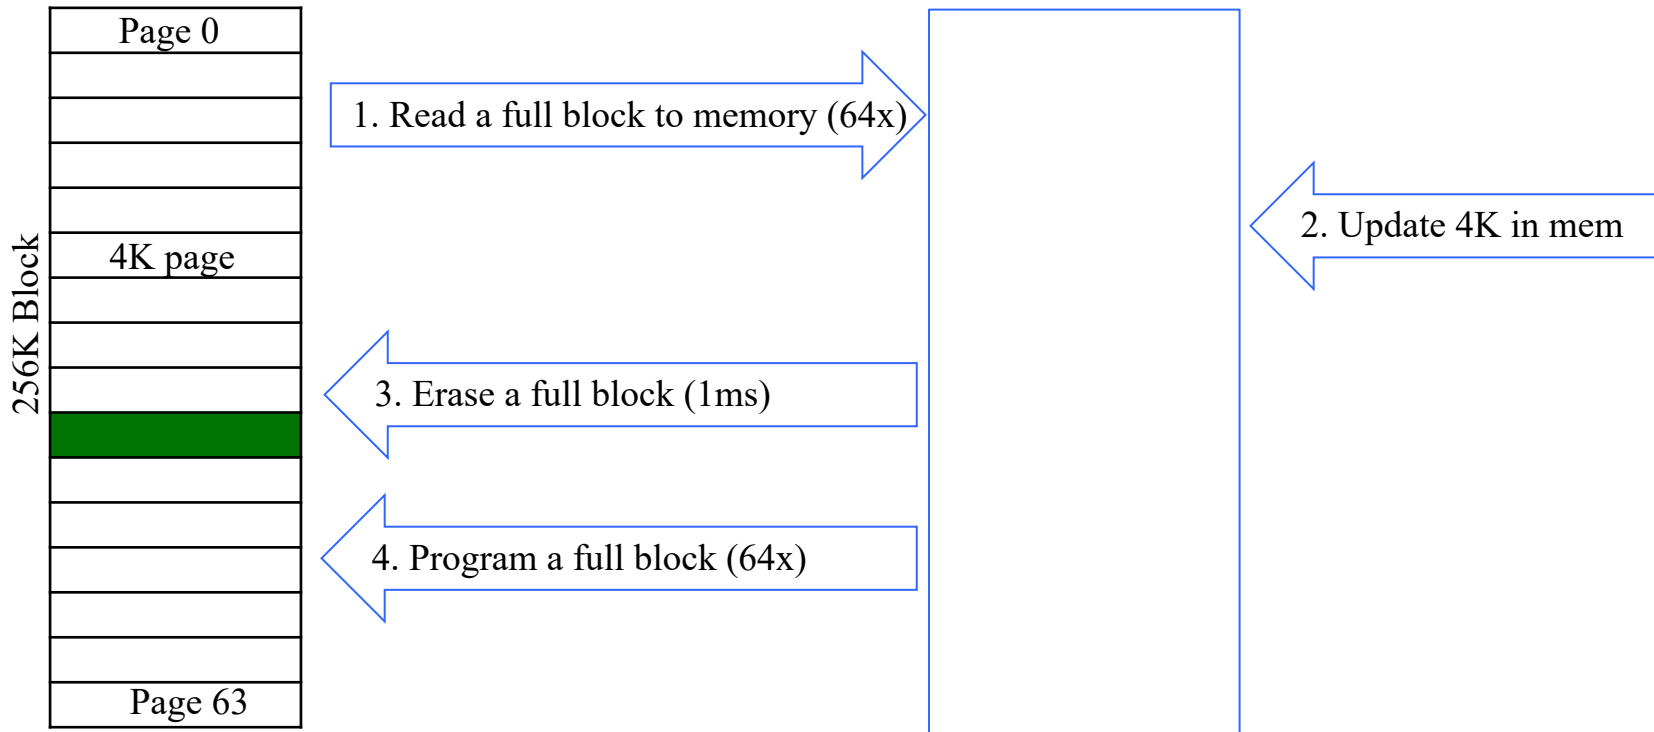
Today

- A lot of concurrency internally
- Placement in internal Flash chips to improve I/O concurrency
 - Spatial parallelism
- Optimize concurrent internal operations
 - Eg spread out mutually blocking operations
- Use lower-cost operations when possible
 - Eg some types of copying/moving faster than others

I/O writes in SSDs

- Let's try to do a small write – where small = 4K
 - Let's say this is an update (not first write)
 - What could we do?
 - Find physical 4K page containing block
 - Cannot write page directly – need to erase block first
 - Read block to memory (read)
 - Update page (modify)
 - Erase block and program full block (write)
 - This is a lot of overhead!
-

I/O writes in SSDs



- Traffic 128x more
- Erase on critical path
- Read-modify-write does not work (write amplification)

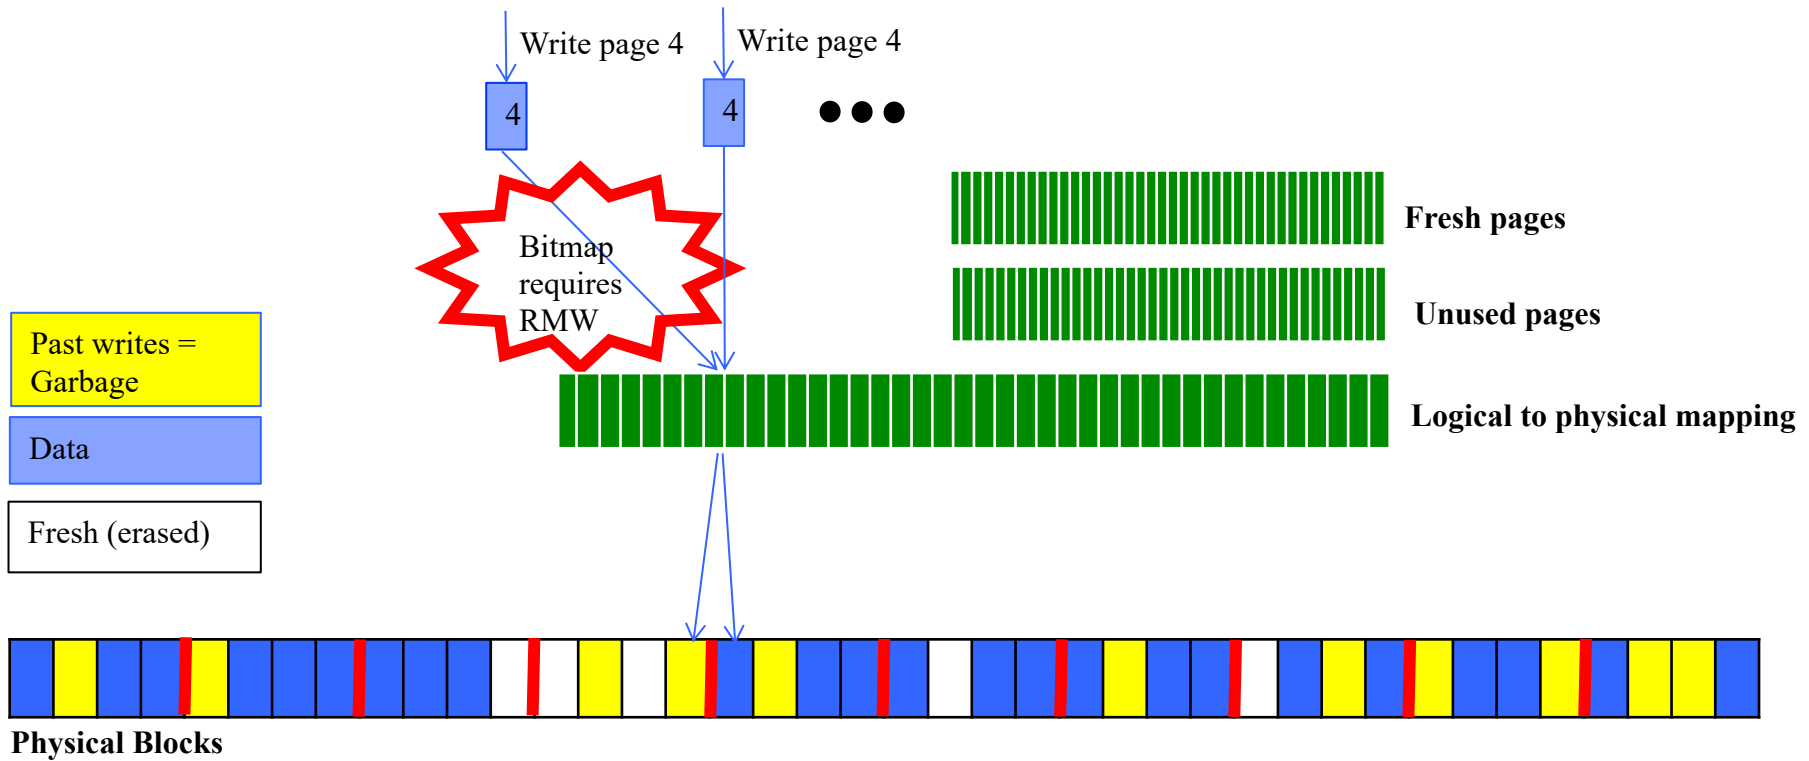
The problem

- Different erase block vs. page sizes
 - E.g. erase block size 256K - erases
 - E.g. page size 4K – updates
- Why? Large blocks => higher density for NAND

Let's rethink

- We cannot overwrite the same page, unless we erase
 - Fine, let's find another page in a block that has been fresh erased
 - Write 4K to this page
 - Remap page to new physical page
 - Free old page
 - Essentially, introduce a logical-physical page mapping
 - First time ever in a storage device for all blocks
 - Requires a lot and non-trivial metadata – like a file system
-


Logical to physical mapping



After a while...

- We end up with:
- Erased blocks that were empty will fill up
- Pages that were re-written will be free leaving holes in various blocks here and there
- At a next write, there will be no location to put the data
 - Cannot write a “hole”, need to erase first
 - Cannot erase a “hole”, need to erase the full block


No “fresh” pages



Garbage Collection

- So at a next write
 - We need to erase a block and create space
 - To erase a block we need to consolidate “holes” in a single erase block
 - This consolidation process is a form of garbage collection
 - Requires
 - Assume a few spare erase blocks
 - Read many pages from various erase blocks
 - Write these pages in fewer erase blocks
 - Update metadata (mappings) and worry about failures
 - All this while performing an I/O write
 - High overhead
 - Unpredictable
-

Create “fresh” pages




Today

- Perform garbage collection asynchronously
 - Reduce overhead of GC
 - Eg segment SSD space between read-mostly and write-mostly data
 - Read-mostly will not be updated so erase blocks do not have “holes”
 - Write-mostly are updated frequently creating many “holes” => empty erase blocks (ready to be erased)
 - Generally, today
 - a complex process
 - high overhead
 - still unpredictable behavior
-

Go back one step

- Can we avoid this problem?
- (1) pages / erase block (# metal contacts) – hurts performance
 - E.g. Erase block size = Page size
 - Erase block size depends on technology and affects density
 - Density critical – probably THE driving force behind market size
- (2) bits / cell (multi-level) – hurts reliability
- (3) page size not as fundamental – important for parallelism
 - No parallelism => expensive read/program
- Maybe there is room for different configurations



Wear Leveling

- Each block can be erased a number of times only
 - After that not possible to update/read
 - If we are unlucky (and naive)
 - One block gets picked all the time to erase
 - After a number of erase cycles it goes bad
 - The SSD needs to reduce size
 - The filesystem (usually) does not handle this
 - Device is not usable
 - What can we do
 - (1) Replace blocks
 - (2) Spread erases across blocks
-

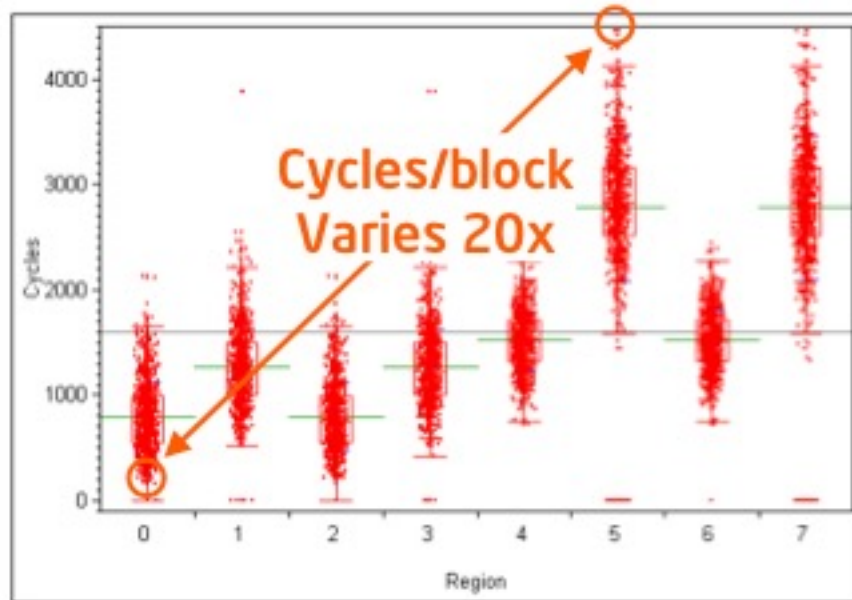
Replace blocks

- No matter what, blocks will go bad at some point
- SSDs have some spare blocks internally (non-negligible percentage)
- When a block goes bad, the SSD detects this
- It remaps logical block to a new physical block – SSD size remains the same
- This is straight-forward (not a common path operation)

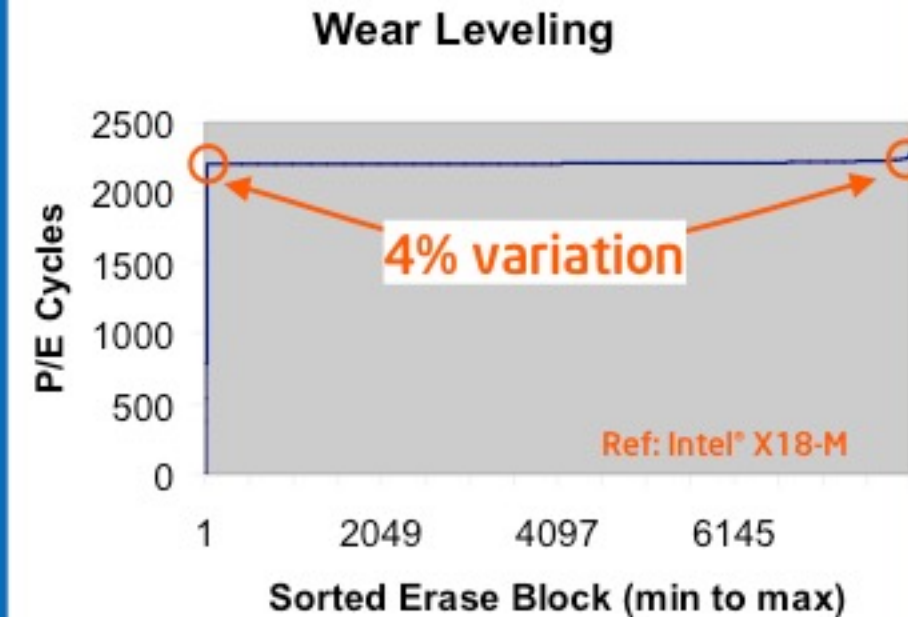
Spread erase operations

- Ensure that each block is erased about the same number of times
 - “Wear-leveling”
- Each block has metadata on how many times it has been erased
- When selecting a block to erase
 - (we said) consult the “fullness” of the block – less full block means less work for GC
 - In addition now consider the erase count
- Elaborate policies
 - E.g. segmentation of space for reads/writes does not help wear leveling

Variability in Wear Leveling



Unsophisticated regioned scheme



More sophisticated scheme

Controllers vary in wear-leveling effectiveness

Poor wear leveling can have high impact

20x in cycles can be 10x or more in RBER

10x in RBER is $10^{\text{ECC}+1}$ in ECC failure rate: 100,000x for 4-bit ECC

Putting it together: SSD Reliability Metrics

SSD UBER values can be $\ll 10^{-15}$

UBER \propto usage: program/erase/read & subsequent retention

Intel® X18-M and X25-M Mainstream SATA SSD (80GB)

- 10 Channels Architecture with 50nm MLC ONFI 1.0 NAND
- 5 years usage, 1000G, 1.2million hrs MTBF
- GB/day client workload @ $1e^{-15}$ UBER $\rightarrow >> 100\text{GB/day}$, 5 years

***Intel® X25-M and X18-M Mainstream SATA SSDs deliver
>5X accepted requirement for clients (20GB/day)***

Intel® X25-E Extreme SATA SSD (32GB)

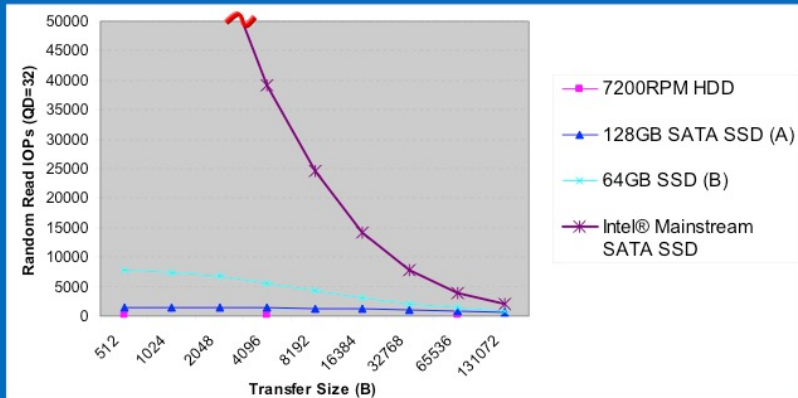
- 10 Channels Architecture with 50nm SLC ONFI 1.0 NAND
- 1000G, 2Million hrs MTBF
- Intel SLC SSD support > 7000 8K 2:1 R/W Random IOPs 24/7, 5 years

***Intel X25-E SLC SSDs support the endurance required
to replace many 15K RPM HDDs for IOPS applications***



SSD Performance


Intel® Mainstream SATA SSD Bridges the HDD Performance Gap Random Read Performance



Performance measurements are made using specific computer systems and/or components and reflect the approximate performance of the technology as measured by those tests. Any difference in system hardware or software design or configuration may affect actual results.



Intel® Mainstream SATA SSD Bridges the HDD Performance Gap (cont'd) Random Write Performance




Performance measurements are made using specific computer systems and/or components and reflect the approximate performance of the technology as measured by those tests. Any difference in system hardware or software design or configuration may affect actual results.



- BUT performance variation a big problem, even today
- More work needed
- SSD controllers and host side become more and more complex

Intel® Mainstream SATA SSDs Save Power: SATA Power Rails With 2 Hour Mobile Workload



What makes a good SSD

- SSD performance not always superior to hard disks (HDDs)
 - Request size matters
 - Request type matters
 - Access pattern matters
 - A HDD raid can be faster than an SSD at the same price
 - A good SSD
 - I/O concurrency
 - Garbage Collection
 - Wear leveling
 - Still room for improvement
-

Outline

- NAND Flash
 - How it works
 - Structure
 - Errors
 - SSDs
 - High concurrency
 - Efficient garbage collection
 - Efficient wear leveling
 - System-level SSD Storage
 - SSD caching
 - Remarks
 - Non-volatile memories (NVM)
-

Background

- Increased need for high-performance storage I/O
 1. Larger file-set sizes \Rightarrow more I/O time
 2. Server virtualization and consolidation \Rightarrow more I/O pressure
- SSDs can mitigate I/O penalties

	SSD	HDD
Throughput (R/W) (MB/s)	277/202	100/90
Response time (ms)	0.17	12.6
IOPS (R/W)	30,000/3,500	150/150
Price/capacity (\$/GB)	\$3	\$0.3
Capacity per device	32 – 120 GB	Up to 3TB

- Mixed SSD and HDD environments are necessary
- Cost-effectiveness: deploy SSDs as HDDs caches

Similar I/O Performance For IOPS Intensive Workload



490 Fiberchannel 15K
RPM drives in 4 racks

Vs.

8 lab prototype SSDs (not
product) internal to server



Unlikely to happen...

- Flash will not replace disks
 - Disk lowest cost/Gbyte
 - Trends not clear – not necessarily in favor of Flash
 - Manufacturing capacity for disks high
 - Would require replacing disk manufacturing capacity with Flash manufacturing capacity – not viable economically
 - More likely
 - Integrate both SSDs and HDDs in storage system
 - E.g. SSD caching
-

Outline

- NAND Flash
 - How it works
 - Structure
 - Errors
 - SSDs
 - Structure
 - Garbage collection
 - Wear leveling
 - System-level SSD Storage
 - SSD caching
 - Remarks
 - Non-volatile memories (NVM)
-

Non-volatile memory (NVM)

- NAND Flash is a type of NVM
 - But
 - Still block addressable (similar to HDDs)
 - Accessible via the same interfaces as HDDs
 - Recent NVMe is FLASH on PCI
 - Other types of NVM are emerging
 - NVM refers to byte-addressable technologies
 - Most prominent
 - Phase change memory (PCM)
 - Magnetic RAM, variant Spin Transfer Torque (STT-RAM)
 - Not clear what will happen, but something will happen
-

Situation 1: More I/O throughput

- Consider a server in the (near) future
 - 1000 cores
 - 1 Gbyte/s/core memory throughput => 1 Tbyte/s
 - 1 Gbit/s/core I/O throughput => 1 Tbit/s
 - Assume all I/O throughput is eventually storage
 - This implies 10,000 spindles/server (assuming 100 MBytes/s/spindle)
 - Or 100,000 spindles/server (assuming 10 MBytes/s/spindle)
 - SSDs/NVMe/NVM is the only way to achieve this compute density
 - Memory-level throughput for persistent storage
-

Situation 2: Less I/O throughput

- Alternative:
 - Let's reduce required I/O by keeping more stuff (data) in memory
 - DRAM is very energy hungry
 - Need to refresh all the time, even if not accessing data
 - Some claim up to 50% of total energy in data-centre
 - Fast persistent storage, and especially NVM (byte-addressable) is currently our best way to address this
 - Practically 0 idle energy consumption
-

The future memory hierarchy

- One memory to rule them all - Unlikely
- NVM to replace DRAM
 - DRAM is faster for writes
 - All NVM memories need to change state in a “permanent” manner for writes and this costs
- NVM/SSD to replace magnetic disks
 - Disk has high areal density which makes it low cost
 - Manufacturing capacity for NVM not adequate to replace disks
- Eventually, complex memory hierarchy
 - DRAM ⇔ NVM ⇔ Flash ⇔ HDDs

How?

- Nobody knows yet ☺
 - Not an easy problem
 - SSDs have been around for almost 10 years now
 - Still not well integrated in storage hierarchy
 - NVM will have issues at multiple layers
 - Devices and packaging
 - What (write) update time will each technology allow?
 - Architectural level
 - What to place in DRAM and what in NVM?
 - Systems software level
 - Can we afford the overhead of the traditional I/O stack?
 - Application level
 - Can we reduce recovery overheads?
 - There is potential for significant impact
-

Questions?

Angelos Bilas
FORTH-ICS

and University of Crete, Greece

bilas@ics.forth.gr
